

Изоляция

Изоляция и системный каталог:

1 терминал

```
access=# create function f() returns integer as $$ select 1;
access=# $$ language sql;
CREATE FUNCTION
```

2 терминал

```
postgres=# \c access
Вы подключены к базе данных "access" как пользователь "postgres".
access=# begin isolation level read committed;
BEGIN
access=# select f();
 f
---
 1
```

1 терминал:

```
access=# create or replace function f() returns integer as $$ select 2;
access=# $$ language sql;
CREATE FUNCTION
access=# select f();
 f
---
 2
```

2 терминал

```
access=# select f();
 f
---
 1
```

1 терминал:

```
access=# create table t(id integer);
CREATE TABLE
```

2 терминал:

```
access=# begin isolation level repeatable read;
BEGIN
```

1 терминал:

```
access=# insert into t values (1);
INSERT 0 1
access=#
```

2 терминал:

```
access=# select * from t;
 id
---
  1
```

1 терминал:

```
access=# insert into t values (2);
INSERT 0 1
access=#
```

2 терминал:

```
access=# select * from t;
 id
----
  1
```

Страницы и версии строк

```
access=# CREATE EXTENSION pageinspect;
ОШИБКА: расширение "pageinspect" уже существует
access=# CREATE TABLE t(s text);
CREATE TABLE
access=# CREATE VIEW t_v AS
access=# SELECT '(0,||lp||)' AS ctid,
access=# CASE lp_flags
access=# WHEN 0 THEN 'unused'
access=# WHEN 1 THEN 'normal'
access=# WHEN 2 THEN 'redirect to '||lp_off
access=# WHEN 3 THEN 'dead'
access=# END AS state,
access=# t_xmin as xmin,
access=# t_xmax as xmax,
access=# CASE WHEN (t_infomask & 256) > 0 THEN 't' END AS xmin_c,
access=# CASE WHEN (t_infomask & 512) > 0 THEN 't' END AS xmin_a,
access=# CASE WHEN (t_infomask & 1024) > 0 THEN 't' END AS xmax_c,
access=# CASE WHEN (t_infomask & 2048) > 0 THEN 't' END AS xmax_a,
access=# t_ctid
access=# FROM heap_page_items(get_raw_page('t',0))
access=# ORDER BY lp;
CREATE VIEW
```

Запускаем транзакцию:

```
access=# begin;
BEGIN
access=# insert into t values ('first');
INSERT 0 1
access=# select txid_current();
 txid_current
-----
        661
```

```
access=# select *, xmin, xmax FROM t;
  s  | xmin | xmax
-----+-----+-----
first |   661 |    0
```

```
access=# select * from t_v;
 ctid | state | xmin | xmax | xmin_c | xmin_a | xmax_c | xmax_a | t_ctid
-----+-----+-----+-----+-----+-----+-----+-----+-----
(0,1) | normal | 661 |    0 |         |         |         |         | (0,1)
```

Создаем точку сохранения:

```
access=# select *, xmin, xmax FROM t;
  s  | xmin | xmax
-----+-----+-----
first |   661 |    0
second | 662 |    0
```

```
access=# savepoint sp;
SAVEPOINT
access=# insert into t values ('second');
INSERT 0 1
access=# select txid_current();
 txid_current
-----
        661
```

```
access=# select * from t_v;
 ctid | state | xmin | xmax | xmin_c | xmin_a | xmax_c | xmax_a | t_ctid
-----+-----+-----+-----+-----+-----+-----+-----+-----
(0,1) | normal | 661 |    0 |         |         |         |         | (0,1)
(0,2) | normal | 662 |    0 |         |         |         |         | (0,2)
```

Делаем откат:

```
access=# rollback to sp;
ROLLBACK
access=# insert into t values ('third');
INSERT 0 1
access=# select txid_current();
 txid_current
-----
        661
```

```
access=# select *, xmin, xmax FROM t;
 s  | xmin | xmax
-----+-----+-----
first |    661 |    0
third |    663 |    0
```

```
access=# select * from t_v;
 ctid | state | xmin | xmax | xmin_c | xmin_a | xmax_c | xmax_a | t_ctid
-----+-----+-----+-----+-----+-----+-----+-----+-----
(0,1) | normal | 661 |    0 | t       |       |       | t       | (0,1)
(0,2) | normal | 662 |    0 |       | t      |       | t       | (0,2)
(0,3) | normal | 663 |    0 | t       |       |       | t       | (0,3)
```

Завершаем транзакцию и смотри на результат:

```
access=# commit;
COMMIT
access=# select *, xmin, xmax FROM t;
 s  | xmin | xmax
-----+-----+-----
first |    661 |    0
third |    663 |    0
```

```
access=# select * from t_v;
 ctid | state | xmin | xmax | xmin_c | xmin_a | xmax_c | xmax_a | t_ctid
-----+-----+-----+-----+-----+-----+-----+-----+-----
(0,1) | normal | 661 |    0 |       |       |       | t       | (0,1)
(0,2) | normal | 662 |    0 |       | t      |       | t       | (0,2)
(0,3) | normal | 663 |    0 |       |       |       | t       | (0,3)
```

Снимки данных

Создаем таблицу и запускаем первую транзакцию:

```
access=# create table t(n integer);
CREATE TABLE
access=# begin;
BEGIN
access=# insert into t(n) values (1);
INSERT 0 1
access=# select txid_current();
 txid_current
-----
        666
```

```
access=# commit;
COMMIT
access=# BEGIN ISOLATION LEVEL REPEATABLE READ;
BEGIN
access=# select * from t;
 n
---
 1
```

```
access=# select txid_current();
 txid_current
-----
        668
```

```
access=# SELECT txid_current_snapshot();
 txid_current_snapshot
-----
668:668:
```

Вторая транзакция:

```
access=# begin isolation level repeatable read;
BEGIN
access=# delete from t;
DELETE 1
access=# select * from t;
n
---
```

```
access=# SELECT txid_current();
txid_current
-----
        669
```

```
access=# SELECT txid_current_snapshot();
txid_current_snapshot
-----
        668:668:
```

Первая транзакция, завершаем обе транзакции:

```
access=# select *, xmin, xmax FROM t;
n | xmin | xmax
---+-----+-----
1 | 667 | 669
```

Создаем функцию и запускаем транзакцию:

```
access=# create function test() returns integer as $$
access$# select pg_sleep(1);
access$# select count(*)::integer from t;
access$# $$ volatile language sql;
CREATE FUNCTION
```

```
access=# select (select count(*) from t) as cnt, test() from generate_series(1,5);
```

Второй терминал:

```
access=# insert into t values(1);
INSERT 0 1
```

Результат в транзакции:

```
access=# select (select count(*) from t) as cnt, test() from generate_series(1,5);
cnt | test
---+-----
0 | 0
0 | 0
0 | 1
0 | 1
0 | 1
```

Запускаем новую транзакцию и повторяем действия:

```
access=# truncate t;
TRUNCATE TABLE
access=# BEGIN ISOLATION LEVEL REPEATABLE READ;
BEGIN
access=# select (select count(*) from t) as cnt, test() from generate_series(1,5);
```

Второй терминал:

```
access=# insert into t values(1);  
INSERT 0 1  
access=#
```

Результат в транзакции:

```
access=# select (select count(*) from t) as cnt, test() from generate_series(1,5);  
cnt | test  
-----+-----  
0   | 0  
0   | 0  
0   | 0  
0   | 0  
0   | 0
```

Запускаем транзакцию:

```
access=# alter function test stable;  
ALTER FUNCTION  
access=# truncate t;  
TRUNCATE TABLE  
access=# BEGIN ISOLATION LEVEL READ committed;  
BEGIN  
access=# select (select count(*) from t) as cnt, test() from generate_series(1,5);  
cnt | test  
-----+-----  
0   | 0  
0   | 0  
0   | 0  
0   | 0  
0   | 0
```

Второй терминал:

```
INSERT 0 1  
access=# insert into t values(1);  
INSERT 0 1  
access=# _
```

Результат транзакции:

```
access=# select (select count(*) from t) as cnt, test() from generate_series(1,5);  
cnt | test  
-----+-----  
0   | 0  
0   | 0  
0   | 0  
0   | 0  
0   | 0
```

Запускаем транзакцию:

```
access=# truncate t;  
TRUNCATE TABLE  
access=# BEGIN ISOLATION LEVEL REPEATABLE READ;  
BEGIN  
access=# select (select count(*) from t) as cnt, test() from generate_series(1,5)
```

На втором терминале:

```
access=# insert into t values(1);  
INSERT 0 1  
access=#
```

Результат транзакции:

```
access=# select (select count(*) from t) as cnt, test() from generate_series(1,5);  
cnt | test  
-----+-----  
0   | 0  
0   | 0  
0   | 0  
0   | 0  
0   | 0
```

НОТ- обновления

Обновления:

```
postgres=# create table t(id integer, s char(2000)) with (fillfactor=75);
CREATE TABLE
postgres=# create index t_id on t(id);
CREATE INDEX
postgres=# create extension pageinspect;
CREATE EXTENSION
postgres=# CREATE VIEW t_v AS
postgres=# SELECT '(0,||lp||)' AS ctid,
postgres=# CASE lp_flags
postgres=# WHEN 0 THEN 'unused'
postgres=# WHEN 1 THEN 'normal'
postgres=# WHEN 2 THEN 'redirect to '||lp_off
postgres=# WHEN 3 THEN 'dead'
postgres=# END AS state,
postgres=# t_xmin || CASE
postgres=# WHEN (t_infomask & 256) > 0 THEN ' (c)'
postgres=# WHEN (t_infomask & 512) > 0 THEN ' (a)'
postgres=# ELSE ''
postgres=# END AS xmin,
postgres=# t_xmax || CASE
postgres=# WHEN (t_infomask & 1024) > 0 THEN ' (c)'
postgres=# WHEN (t_infomask & 2048) > 0 THEN ' (a)'
postgres=# ELSE ''
postgres=# END AS xmax,
postgres=# CASE WHEN (t_infomask2 & 16384) > 0 THEN 't' END AS hhu,
postgres=# CASE WHEN (t_infomask2 & 32768) > 0 THEN 't' END AS hot,
postgres=# t_ctid
postgres=# FROM heap_page_items(get_raw_page('t',0))
postgres=# ORDER BY lp;
CREATE VIEW
postgres=# create view t_id_v as
postgres=# select itemoffset,
postgres=# ctid
postgres=# from bt_page_items('t_id',1);
CREATE VIEW
```

```
postgres=# insert into t(s) values ('A');
INSERT 0 1
postgres=# update t set s='B';
UPDATE 1
postgres=# update t set s='C';
UPDATE 1
postgres=# update t set s='D';
UPDATE 1
postgres=# select * from t_id_v;
itemoffset | ctid
-----+-----
1 | (0,1)
(1 row)

postgres=# select * from t_v;
ctid | state | xmin | xmax | hhu | hot | t_ctid
-----+-----+-----+-----+-----+-----+-----
(0,1) | normal | 573 (c) | 574 (c) | t | | (0,2)
(0,2) | normal | 574 (c) | 575 (c) | t | t | (0,3)
(0,3) | normal | 575 (c) | 576 | t | t | (0,4)
(0,4) | normal | 576 | 0 (a) | | t | (0,4)
(4 rows)
```

Очистка:

```
postgres=# update t set s='E';
UPDATE 1
postgres=# select * from t_v;
ctid | state | xmin | xmax | hhu | hot | t_ctid
-----+-----+-----+-----+-----+-----+-----
(0,1) | redirect to 4 | 577 | 0 (a) | | t | (0,2)
(0,2) | normal | 577 (c) | 578 (c) | t | t | (0,3)
(0,3) | normal | 578 (c) | 579 | t | t | (0,5)
(0,4) | normal | 576 (c) | 577 (c) | t | t | (0,2)
(0,5) | normal | 579 | 0 (a) | | t | (0,5)
(5 rows)

postgres=# update t set s='F';
UPDATE 1
postgres=# update t set s='G';
UPDATE 1
postgres=# select * from t_v;
ctid | state | xmin | xmax | hhu | hot | t_ctid
-----+-----+-----+-----+-----+-----+-----
(0,1) | redirect to 4 | 577 (c) | 578 (c) | t | t | (0,3)
(0,2) | normal | 578 (c) | 579 | t | t | (0,5)
(0,3) | normal | 576 (c) | 577 (c) | t | t | (0,2)
(0,4) | normal | 579 | 0 (a) | | t | (0,5)
(5 rows)

postgres=# update t set s='H';
UPDATE 1
postgres=# select * from t_v;
ctid | state | xmin | xmax | hhu | hot | t_ctid
-----+-----+-----+-----+-----+-----+-----
(0,1) | redirect to 5 | 580 | 0 (a) | | t | (0,2)
(0,2) | normal | 580 (c) | 581 (c) | t | t | (0,3)
(0,3) | normal | 581 (c) | 582 (c) | t | t | (0,4)
(0,4) | normal | 582 (c) | 583 (c) | t | t | (0,5)
(0,5) | normal | 583 (c) | 584 | t | t | (1,1)
(5 rows)
```

Разрыв hot- цепочки:

Запускаем транзакцию

```
postgres=# begin isolation level repeatable read;
BEGIN
postgres=# select count(*) from t;
count
-----
1
(1 row)
```

Открываем новый терминал:

```
postgres=# update t set s='I';
UPDATE 1
postgres=# update t set s='J';
UPDATE 1
postgres=# update t set s='K';
UPDATE 1
postgres=# select * from t_v;
ctid | state | xmin | xmax | hhu | hot | t_ctid
-----+-----+-----+-----+-----+-----+-----
(0,1) | redirect to 2 | 580 (c) | 581 (c) | t | t | (0,3)
(0,2) | normal | 581 (c) | 582 (c) | t | t | (0,4)
(0,3) | normal | 582 (c) | 583 (c) | t | t | (0,5)
(0,4) | normal | 583 (c) | 584 | t | t | (1,1)
(5 rows)

postgres=# update t set s='L';
UPDATE 1
postgres=# select * from t_v;
ctid | state | xmin | xmax | hhu | hot | t_ctid
-----+-----+-----+-----+-----+-----+-----
(0,1) | redirect to 2 | 580 (c) | 581 (c) | t | t | (0,3)
(0,2) | normal | 581 (c) | 582 (c) | t | t | (0,4)
(0,3) | normal | 582 (c) | 583 (c) | t | t | (0,5)
(0,4) | normal | 583 (c) | 584 | t | t | (1,1)
(5 rows)

postgres=# select * from t_id_v;
itemoffset | ctid
-----+-----
1 | (0,1)
2 | (1,1)
(2 rows)
```

Очистка

Создаем таблицу:

```
access=# CREATE TABLE t(id integer) WITH (autovacuum_enabled = off);
CREATE TABLE
access=# INSERT INTO t SELECT gen.id FROM generate_series(1,1000000) gen(id);
INSERT 0 1000000
access=# CREATE INDEX t_id ON t(id);
CREATE INDEX
access=# ALTER SYSTEM SET maintenance_work_mem = '1MB';
ALTER SYSTEM
access=# select pg_reload_conf();
pg_reload_conf
-----
t
```

Обновляем строки:

```
access=# UPDATE t SET id = id + 1;
UPDATE 1000000
```

Запускаем очистку и транзакцию:

```
access=# VACUUM VERBOSE t;
ИНФОРМАЦИЯ: очистка "public.t"
ИНФОРМАЦИЯ: просканирован индекс "t_id", удалено версий строк: 174472
ПОДРОБНОСТИ: CPU: пользов.: 0.14 с, система: 0.00 с, прошло: 0.14 с
ИНФОРМАЦИЯ: "t": удалено версий строк: 174472, обработано страниц: 772
ПОДРОБНОСТИ: CPU: пользов.: 0.00 с, система: 0.00 с, прошло: 0.00 с
ИНФОРМАЦИЯ: просканирован индекс "t_id", удалено версий строк: 174472
ПОДРОБНОСТИ: CPU: пользов.: 0.14 с, система: 0.00 с, прошло: 0.13 с
ИНФОРМАЦИЯ: "t": удалено версий строк: 174472, обработано страниц: 772
ПОДРОБНОСТИ: CPU: пользов.: 0.00 с, система: 0.00 с, прошло: 0.00 с
ИНФОРМАЦИЯ: просканирован индекс "t_id", удалено версий строк: 174472
ПОДРОБНОСТИ: CPU: пользов.: 0.12 с, система: 0.00 с, прошло: 0.12 с
```

```
access=# SELECT * FROM pg_stat_progress_vacuum \gx
(0 строк)
```

Восстанавливаем значения измененного параметра:

```
access=# SELECT pg_reload_conf();
pg_reload_conf
-----
t
```

Смотрим размер файла:

```
access=# SELECT pg_size_pretty(pg_table_size('t'));
pg_size_pretty
-----
69 MB
```

```
access=# DELETE FROM t WHERE random() < 0.9;
DELETE 899579
access=# VACUUM t;
VACUUM
access=# SELECT pg_size_pretty(pg_table_size('t'));
pg_size_pretty
-----
69 MB
```

Полная очистка:

```
access=# TRUNCATE t;
TRUNCATE TABLE
access=# INSERT INTO t SELECT gen.id FROM generate_series(1,1000000) gen(id);
access=# ;
ОШИБКА: ошибка синтаксиса (примерное положение: "|" )
СТРОКА 1: ...INTO t SELECT gen.id FROM generate_series(1,1000000) gen(id)|
^
access=# INSERT INTO t SELECT gen.id FROM generate_series(1,1000000) gen(id);
INSERT 0 1000000
access=# SELECT pg_size_pretty(pg_table_size('t'));
pg_size_pretty
-----
35 MB
```

```
access=# DELETE FROM t WHERE random() < 0.9;
DELETE 900081
access=# VACUUM FULL t;
VACUUM
access=# SELECT pg_size_pretty(pg_table_size('t'));
pg_size_pretty
-----
3544 kB
```

Автоочистка

Настраиваем очистку:

```
access=# alter system set autovacuum_vacuum_threshold = 0;
ALTER SYSTEM
access=# alter system set autovacuum_vacuum_scale_factor = 0.1;
ALTER SYSTEM
access=# alter system set autovacuum_naptime = '1s';
ALTER SYSTEM
access=# select pg_reload_conf();
pg_reload_conf
-----
t
```

Создаем и заполняем таблицу:

```
access=# create table t(n integer);
CREATE TABLE
access=# insert into t(n) select 1 from generate_series(1,100000);
INSERT 0 100000
access=#
```

```
access=# select pg_size_pretty(pg_table_size('t'));
pg_size_pretty
-----
3568 kB
```

Делаем обновление таблицы

```
access=# update t set n=n+1 where random()<0.055;
UPDATE 5503
access=# update t set n=n+1 where random()<0.055;
UPDATE 5485
access=# update t set n=n+1 where random()<0.055;
UPDATE 5539
access=# update t set n=n+1 where random()<0.055;
UPDATE 5599
access=# update t set n=n+1 where random()<0.055;
UPDATE 5535
access=# update t set n=n+1 where random()<0.055;
UPDATE 5487
access=# update t set n=n+1 where random()<0.055;
UPDATE 5524
access=# update t set n=n+1 where random()<0.055;
UPDATE 5470
access=# update t set n=n+1 where random()<0.055;
UPDATE 5449
access=# update t set n=n+1 where random()<0.055;
UPDATE 5382
access=# update t set n=n+1 where random()<0.055;
UPDATE 5545
access=# update t set n=n+1 where random()<0.055;
UPDATE 5445
access=# update t set n=n+1 where random()<0.055;
UPDATE 5379
access=# update t set n=n+1 where random()<0.055;
UPDATE 5438
access=# update t set n=n+1 where random()<0.055;
UPDATE 5433
access=# update t set n=n+1 where random()<0.055;
UPDATE 5580
access=# update t set n=n+1 where random()<0.055;
UPDATE 5507
access=# update t set n=n+1 where random()<0.055;
UPDATE 5509
access=# update t set n=n+1 where random()<0.055;
UPDATE 5441
access=# update t set n=n+1 where random()<0.055;
```

Оценка разности

```
access=# SELECT * FROM pg_stat_all_tables WHERE relname='t' \gx
-[ RECORD 1 ]-----+-----
relid              | 16540
schemaname         | public
relname            | t
seq_scan           | 20
seq_tup_read       | 2000000
idx_scan           | 0
idx_tup_fetch      | 0
n_tup_ins          | 100000
n_tup_upd          | 100709
n_tup_del          | 0
n_tup_hot_upd      | 78509
n_live_tup         | 100000
n_dead_tup         | 5459
n_mod_since_analyze | 0
last_vacuum        | 
last_autovacuum    | 2021-10-28 15:02:42.54631+03
last_analyze       | 
last_autoanalyze   | 2021-10-28 15:02:45.597224+03
vacuum_count       | 0
autovacuum_count   | 4
analyze_count      | 0
autoanalyze_count  | 11

access=# SELECT pg_size_pretty(pg_table_size('t'));
pg_size_pretty
-----
4416 kB
```


Заморозка

Запуск транзакции:

```
access=# begin;
BEGIN
access=# create table t(n integer);
CREATE TABLE
access=# copy t from stdin with freeze;
Вводите данные для копирования, разделяя строки переводом строки.
Закончите ввод строкой '\.' или сигналом EOF.
>> 1
>> 2
>> 3
>> \.
COPY 3
access=# commit;
COMMIT
```

```
access=# CREATE VIEW t_vs AS
access=# SELECT '(0,||lp||)' as ctid,
access=# CASE lp_flags
access=# WHEN 0 THEN 'unused'
access=# WHEN 1 THEN 'normal'
access=# WHEN 3 THEN 'dead'
access=# END AS state,
access=# t_xmin AS xmin,
access=# age(t_xmin) AS xmin_age,
access=# CASE WHEN (t_infomask & 256) > 0 THEN 't' END AS xmin_c,
access=# CASE WHEN (t_infomask & 512) > 0 THEN 't' END AS xmin_a,
access=# t_xmax AS xmax,
access=# t_ctid
access=# FROM heap_page_items(get_raw_page('t',0))
access=# ORDER BY lp;
CREATE VIEW
access=# select * from t_vs;
   ctid | state | xmin | xmin_age | xmin_c | xmin_a | xmax | t_ctid
-----+-----+-----+-----+-----+-----+-----+-----
(0,1) | normal | 727 | 2 | t | t | 0 | (0,1)
(0,2) | normal | 727 | 2 | t | t | 0 | (0,2)
(0,3) | normal | 727 | 2 | t | t | 0 | (0,3)
```

Изоляция:

Транзакция на втором терминале

```
access=# begin isolation level repeatable read;
BEGIN
access=# select txid_current();
txid_current
-----
729
```

Запускаем транзакцию на первом терминале:

```
access=# begin;
BEGIN
access=# truncate t;
TRUNCATE TABLE
access=# copy t from stdin with freeze;
Вводите данные для копирования, разделяя строки переводом строки.
Закончите ввод строкой '\.' или сигналом EOF.
>> 10
>> 20
>> 30
>> \.
COPY 3
access=# commit;
COMMIT
```

На второй транзакции:

```
access=# select * from t;
n
-----
10
20
30
```

```
access=# ALTER SYSTEM SET autovacuum = off;
ALTER SYSTEM
access=# ALTER SYSTEM SET vacuum_freeze_min_age = 1000;
ALTER SYSTEM
access=# ALTER SYSTEM SET vacuum_freeze_table_age = 10000;
ALTER SYSTEM
access=# ALTER SYSTEM SET autovacuum_freeze_max_age = 100000;
ALTER SYSTEM
access=# SELECT datname, datfrozenxid, age(datfrozenxid) FROM pg_database;
datname | datfrozenxid | age
-----+-----+-----
postgres | 562 | 169
access | 562 | 169
template0 | 562 | 169
template1 | 562 | 169
data_lowlevel | 562 | 169

access=# ;
ОШИБКА: отношение "pg_database|" не существует
СТРОКА 1: ...ECT datname, datfrozenxid, age(datfrozenxid) FROM pg_databas.
..
^
access=# SELECT datname, datfrozenxid, age(datfrozenxid) FROM pg_database;
   datname | datfrozenxid | age
-----+-----+-----
postgres | 562 | 169
access | 562 | 169
template0 | 562 | 169
template1 | 562 | 169
data_lowlevel | 562 | 169
```

```
access=# VACUUM t;
VACUUM
access=# SELECT datname, datfrozenxid, age(datfrozenxid) FROM pg_database;
   datname | datfrozenxid | age
-----+-----+-----
postgres | 562 | 169
access | 562 | 169
template0 | 562 | 169
template1 | 562 | 169
data_lowlevel | 562 | 169
```

Буферный кэш

Создаем, заполняем таблицу и смотрим размеры:

```
-----
postgres=# CREATE TABLE t(n integer);
CREATE TABLE
postgres=# INSERT INTO t SELECT 1 FROM generate_series(1,10000);
INSERT 0 10000
postgres=# SELECT setting FROM pg_settings WHERE name = 'block_size';
setting
-----
8192
(1 row)

postgres=# SELECT pg_table_size('t') / 8192;
?column?
-----
48
(1 row)

postgres=# SELECT pg_relation_size('t','main') / 8192;
?column?
-----
45
(1 row)

postgres=# SELECT pg_relation_size('t','fsm') / 8192;
?column?
-----
3
(1 row)

postgres=# CREATE EXTENSION pg_buffercache;
CREATE EXTENSION
postgres=# SELECT CASE relforknumber
postgres-#         WHEN 0 THEN 'main'
postgres-#         WHEN 1 THEN 'fsm'
postgres-#         WHEN 2 THEN 'vm'
postgres-#         END relfork,
postgres-#         count(*)
postgres-# FROM   pg_buffercache b,
postgres-#         pg_class c
postgres-# WHERE  b.reldatabase = (
postgres-#         SELECT oid FROM pg_database WHERE datname = current_database
postgres-#         )
postgres-# AND   c.oid = 't'::regclass
postgres-# AND   b.relfilenode = c.relfilenode
postgres-# GROUP BY 1;
 relfork | count
-----+-----
fsm      |      2
main     |     45
(2 rows)
```

Количество грязных буферов:

```
postgres=# select count(*) from pg_buffercache b where isdirty;
count
-----
0
(1 row)
```

Контрольная строка:

```
postgres=# select count(*) from pg_buffercache b where isdirty;
count
-----
0
(1 row)
```

Журнал предзаписи

```
postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8865B08
(1 row)

postgres=# CREATE TABLE t(id integer PRIMARY KEY);
CREATE TABLE
postgres=# INSERT INTO t VALUES (1),(2),(3);
INSERT 0 3
postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8878B70
(1 row)

postgres=# SELECT '0/2C480D70'::pg_lsn - '0/2C46450C'::pg_lsn;
?column?
-----
116836
(1 row)

postgres=# SELECT pg_walfile_name('0/2C46450C');
pg_walfile_name
-----
000000010000000000000002C
(1 row)
```

Контрольная точка

Настройки контрольной точки:

```
postgres=# ALTER SYSTEM SET checkpoint_timeout = '30s';
ALTER SYSTEM
postgres=# ALTER SYSTEM SET min_wal_size = '16MB';
ALTER SYSTEM
postgres=# ALTER SYSTEM SET max_wal_size = '16MB';
ALTER SYSTEM
postgres=# SELECT pg_reload_conf();
pg_reload_conf
-----
t
(1 row)
```

Нагрузка:

```
postgres=# SELECT pg_stat_reset_shared('bgwriter');
pg_stat_reset_shared
-----
(1 row)

postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8878C58
(1 row)

postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8878C58
(1 row)

postgres=# SELECT pg_size_pretty('0/327F3210'::pg_lsn - '0/2D0B6BF8'::pg_lsn);
pg_size_pretty
-----
87 MB
(1 row)

postgres=# SELECT checkpoints_timed, checkpoints_req FROM pg_stat_bgwriter;
checkpoints_timed | checkpoints_req
-----+-----
2 | 0
(1 row)
```

Возврат настроек:

```
postgres=# ALTER SYSTEM RESET ALL;
ALTER SYSTEM
postgres=# SELECT pg_reload_conf();
pg_reload_conf
-----
t
(1 row)
```

Настройка журнала

ON:

```
postgres=# ALTER SYSTEM SET full_page_writes = on;
ALTER SYSTEM
postgres=# ALTER SYSTEM SET wal_compression = on;
ALTER SYSTEM
postgres=# SELECT pg_reload_conf();
pg_reload_conf
-----
t
(1 row)

postgres=# checkpoint;
CHECKPOINT
postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8878F50
(1 row)

postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8878F50
(1 row)

postgres=# SELECT pg_size_pretty('0/3623DD38'::pg_lsn - '0/35A67878'::pg_lsn);
pg_size_pretty
-----
8025 kB
(1 row)
```

OFF:

```
postgres=# ALTER SYSTEM SET full_page_writes = off;
ALTER SYSTEM
postgres=# SELECT pg_reload_conf();
pg_reload_conf
-----
t
(1 row)

postgres=# checkpoint;
CHECKPOINT
postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8878E48
(1 row)

postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8878E48
(1 row)

postgres=# SELECT pg_size_pretty('0/35A677C0'::pg_lsn - '0/347BC348'::pg_lsn);
pg_size_pretty
-----
19 MB
(1 row)

postgres=# SHOW data_checksums;
data_checksums
-----
off
(1 row)
```

Сжатие:

```
postgres=# SHOW full_page_writes;
full_page_writes
-----
on
(1 row)

postgres=# checkpoint;
CHECKPOINT
postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8878D40
(1 row)

postgres=# SELECT pg_current_wal_insert_lsn();
pg_current_wal_insert_lsn
-----
0/8878D40
(1 row)

postgres=# SELECT pg_size_pretty('0/347BC290'::pg_lsn - '0/33429A54'::pg_lsn);
pg_size_pretty
-----
20 MB
(1 row)
```

Блокировки объектов

Создаем базу данных, таблицу и заносим данные:

```
postgres=# create database locks_obj;
CREATE DATABASE
postgres=# \c locks_obj
Вы подключены к базе данных "locks_obj" как пользователь "postgres".
locks_obj=# create table accounts(acc_no integer primary key, amount);
ОШИБКА:  ошибка синтаксиса (примерное положение: ")")
СТРОКА 1: create table accounts(acc_no integer primary key, amount);
                                ^
locks_obj=# create table accounts(acc_no integer primary key, amount numeric);
CREATE TABLE
locks_obj=# insert into accounts values (1,1000.00), (2, 2000.00), (3, 3000.00);
INSERT 0 3
```

Начинаем транзакцию, выводим одну строку:

```
locks_obj=# \c locks_obj
Вы подключены к базе данных "locks_obj" как пользователь "postgres".
locks_obj=# select pg_backend_pid();
pg_backend_pid
-----
345996
(1 строка)
```

```
locks_obj=# begin;
BEGIN
locks_obj=# select * from accounts where acc_no=1;
 acc_no | amount 
-----+-----
1 | 1000.00
(1 строка)
```

Блокировка вне транзакции:

```
locks_obj=# select locktype, relation::Regclass, virtualxid as virtxid, transactionid as xid, mode, granted from pg_locks where pid = 345996;
 locktype | relation | virtxid | xid | mode | granted 
-----+-----+-----+----+-----+-----
relation | pg_locks | 4/7117 |    | AccessShareLock | t
virtualxid |          | 4/7117 |    | ExclusiveLock | t
(2 строки)
```

Блокировка в транзакции:

```
locks_obj=# begin;
BEGIN
locks_obj=# select * from accounts where acc_no=1;
 acc_no | amount 
-----+-----
1 | 1000.00
(1 строка)

locks_obj=#
locks_obj=# select locktype, relation::Regclass, virtualxid as virtxid, transactionid as xid, mode, granted from pg_locks where pid = 345996;
 locktype | relation | virtxid | xid | mode | granted 
-----+-----+-----+----+-----+-----
relation | pg_locks |          |    | AccessShareLock | t
relation | accounts_pkey |          |    | AccessShareLock | t
relation | accounts |          |    | AccessShareLock | t
virtualxid |          | 4/7119 |    | ExclusiveLock | t
(4 строки)
```

Выполняем транзакцию, читаем одну строку, блокируем:

```
locks_obj=# begin;
BEGIN
locks_obj=# select * from accounts where acc_no=1;
 acc_no | amount 
-----+-----
1 | 1000.00
(1 строка)

locks_obj=#
locks_obj=# select locktype, relation::Regclass, virtualxid as virtxid, transactionid as xid, mode, granted from pg_locks where pid = 345996;
 locktype | relation | virtxid | xid | mode | granted 
-----+-----+-----+----+-----+-----
relation | pg_locks |          |    | AccessShareLock | t
relation | accounts_pkey |          |    | AccessShareLock | t
relation | accounts |          |    | AccessShareLock | t
virtualxid |          | 4/7119 |    | ExclusiveLock | t
(4 строки)
```


Изменяем параметры:

```
locks_obj=# alter system set log_lock_waits=on;
ALTER SYSTEM
locks_obj=# alter system set deadlock_timeout='100ms';
ALTER SYSTEM
locks_obj=# select pg_reload_conf();
pg_reload_conf
-----
t
(1 строка)
```

Первый терминал:

```
locks_obj=# begin;
BEGIN
locks_obj=# update accounts set amount=10.00 where acc_no=1;
UPDATE 1
```

Второй терминал:

```
locks_obj=# begin;
BEGIN
locks_obj=# update accounts set amount=100.00 where acc_no=1;
```

Первый терминал:

```
locks_obj=# SELECT PG_SLEEP(1);
pg_sleep
-----
(1 строка)

locks_obj=# COMMIT;
COMMIT
```

На втором терминале завершаем транзакцию.

Блокировки строк

Создаем
представление:

```
locks_obj=# CREATE VIEW locks AS SELECT pid, locktype, CASE locktype WHEN 'relat
ion' THEN relation::REGCLASS::text WHEN 'virtualxid' THEN virtualxid::text WHEN
'transactionid' THEN transactionid::text WHEN 'tuple' THEN relation::REGCLASS::t
ext||':'||tuple::text END AS lockid, mode, granted FROM pg_locks;
CREATE VIEW
```

Первый
терминал:

```
locks_obj=# begin;
BEGIN
locks_obj=# select txid_current(), pg_backend_pid();
txid_current | pg_backend_pid
-----+-----
606 | 356219
(1 строка)

locks_obj=# update accounts set amount=amount+100.00 where acc_no=1;
UPDATE 1
```

Второй
терминал:

```
locks_obj=# begin;
BEGIN
locks_obj=# select txid_current(), pg_backend_pid();
txid_current | pg_backend_pid
-----+-----
607 | 356862
(1 строка)

locks_obj=# UPDATE accounts SET amount = amount + 100.00 WHERE acc_no = 1;
```

Третий
терминал:

```
locks_obj=# begin;
BEGIN
locks_obj=# SELECT txid_current(), pg_backend_pid();
txid_current | pg_backend_pid
-----+-----
608 | 357061
(1 строка)

locks_obj=# UPDATE accounts SET amount = amount + 100.00 WHERE acc_no = 1;
```

Блокировка для первой транзакции:

```
locks_obj=# SELECT * FROM locks WHERE pid =356219;
```

pid	locktype	lockid	mode	granted
356219	relation	accounts_pkey	RowExclusiveLock	t
356219	relation	accounts	RowExclusiveLock	t
356219	virtualxid	6/1456	ExclusiveLock	t
356219	tuple	accounts:9	ExclusiveLock	t
356219	transactionid	609	ExclusiveLock	t
356219	transactionid	608	ShareLock	f

(6 строк)

Блокировка для второй транзакции:

```
locks_obj=# SELECT * FROM locks WHERE pid =356862;
```

pid	locktype	lockid	mode	granted
356862	relation	accounts_pkey	RowExclusiveLock	t
356862	relation	accounts	RowExclusiveLock	t
356862	virtualxid	8/143	ExclusiveLock	t
356862	transactionid	610	ExclusiveLock	t
356862	tuple	accounts:9	ExclusiveLock	f

(5 строк)

Блокировка для третьей:

pid	locktype	lockid	mode	granted
357061	relation	pg_stat_activity	AccessShareLock	t
357061	relation	pg_locks	AccessShareLock	t
357061	relation	locks	AccessShareLock	t
357061	relation	accounts_pkey	RowExclusiveLock	t
357061	relation	accounts	RowExclusiveLock	t
357061	virtualxid	9/21	ExclusiveLock	t
357061	relation	pg_database_oid_index	AccessShareLock	t
357061	relation	pg_authid_oid_index	AccessShareLock	t
357061	transactionid	608	ExclusiveLock	t
357061	relation	pg_authid_rolname_index	AccessShareLock	t
357061	relation	pg_database	AccessShareLock	t
357061	relation	pg_authid	AccessShareLock	t
357061	relation	pg_database_datname_index	AccessShareLock	t

(13 строк)

Посмотрим блокирующие процессы:

```
locks_obj=# SELECT pid, wait_event_type, wait_event, pg_blocking_pids(pid) FROM pg_stat_activity WHERE backend_type = 'client backend';
```

pid	wait_event_type	wait_event	pg_blocking_pids
356366	Client	ClientRead	{}
356200			{}
133337	Client	ClientRead	{}
356219	Lock	transactionid	{357061}
356862	Lock	tuple	{356219}
357061			{}

(6 строк)

Взаимодействие трех транзакций

Первая транзакция:

```
locks_obj=# begin;
BEGIN
locks_obj=# UPDATE accounts SET amount = amount - 100.00 WHERE acc_no = 1;
UPDATE 1
locks_obj=#
```

Вторая транзакция:

```
locks_obj=# begin;
BEGIN
locks_obj=# UPDATE accounts SET amount = amount - 100.00 WHERE acc_no = 2;
UPDATE 1
locks_obj=#
```

Третья транзакция:

```
locks_obj=# begin;
BEGIN
locks_obj=# UPDATE accounts SET amount = amount - 100.00 WHERE acc_no = 3;
UPDATE 1
locks_obj=#
```

Первая транзакция:

```
locks_obj=## UPDATE accounts SET amount = amount + 100.00 WHERE acc_no = 2;
```

Вторая транзакция:

```
locks_obj=## UPDATE accounts SET amount = amount + 100.00 WHERE acc_no = 3;
```

Третья транзакция, обновление строки на 1 и 2 транзакции завершатся:

```
locks_obj=## UPDATE accounts SET amount = amount + 100.00 WHERE acc_no = 1;
ОШИБКА: обнаружена взаимоблокировка
ПОДРОБНОСТИ: Процесс 357061 ожидает в режиме ShareLock блокировку "транзакция 614"; заблокирован процессом 356219.
Процесс 356219 ожидает в режиме ShareLock блокировку "транзакция 615"; заблокирован процессом 356862.
Процесс 356862 ожидает в режиме ShareLock блокировку "транзакция 616"; заблокирован процессом 357061.
ПОДСКАЗКА: Подробности запроса смотрите в протоколе сервера.
КОНТЕКСТ: при изменении кортежа (0,11) в отношении "accounts"
```

Взаимоблокировка двух операций UPDATE

Первая транзакция:

```
locks_obj=# begin;
BEGIN
locks_obj=## DECLARE c1 CURSOR FOR SELECT * FROM accounts ORDER BY acc_no FOR UPDATE;
DECLARE CURSOR
```

Вторая транзакция:

```
locks_obj=# begin;
BEGIN
locks_obj=## DECLARE c2 CURSOR FOR SELECT * FROM accounts ORDER BY acc_no DESC FOR UPDATE;
DECLARE CURSOR
```

Первая транзакция:

```
locks_obj=## FETCH c1;
acc_no | amount
-----+-----
      1 | 300.00
(1 строка)
```

Вторая транзакция:

```
locks_obj=## FETCH c2;
acc_no | amount
-----+-----
      3 | 3000.00
(1 строка)
locks_obj=##
```

Первая транзакция:

```
locks_obj=## FETCH c1;
acc_no | amount
-----+-----
      2 | 2000.00
(1 строка)
```


Вторая транзакция (задержка):

```
locks_obj=##  
locks_obj=## FETCH c2;
```

Первая транзакция:

```
locks_obj=## FETCH c1;  
ОШИБКА: обнаружена взаимоблокировка  
ПОДРОБНОСТИ: Процесс 356219 ожидает в режиме ShareLock блокировку "транзакция 6  
18"; заблокирован процессом 356862.  
Процесс 356862 ожидает в режиме ShareLock блокировку "транзакция 617"; заблокиро  
ван процессом 356219.  
ПОДСКАЗКА: Подробности запроса смотрите в протоколе сервера.  
КОНТЕКСТ: при блокировке кортежа (0,3) в отношении "accounts"  
locks_obj=1#
```

Появление на второй:

```
acc_no | amount  
-----+-----  
      3 | 3000.00  
(1 строка)  
locks_obj=##
```